

# Database Systems

Instructor: Bilal Khalid Dar



# Relational Algebra

# DBMS Architecture

How does a SQL engine work ?

- SQL query → relational algebra plan
- Relational algebra plan → Optimized plan
- Execute each operator of the plan

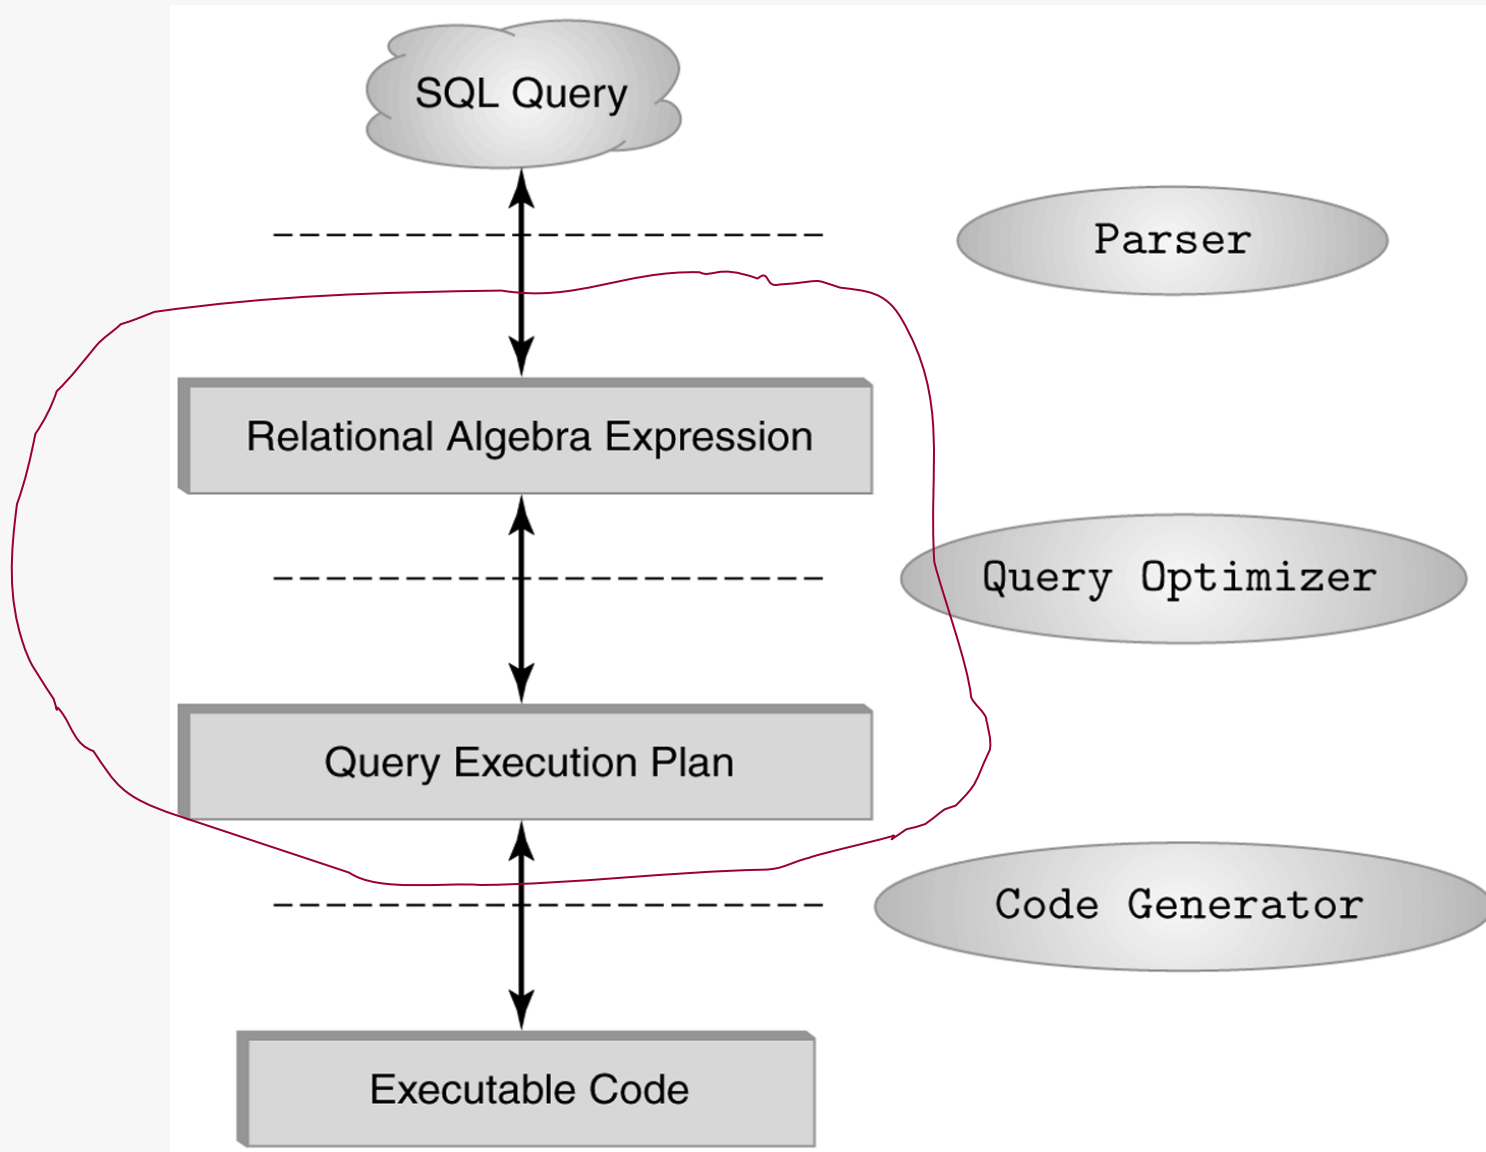
# What is an Algebra?

- A language based on operators and a domain of values
- Operators map values taken from the domain into other domain values
- Hence, an expression involving operators and arguments produces a value in the domain
- When the domain is a set of all relations (and the operators are as described later), we get the *relational algebra*
- We refer to the expression as a *query* and the value produced as the *query result*

# Relational Algebra

- Five operators:
  - Union:  $\cup$
  - Difference:  $-$
  - Selection:  $\sigma$
  - Projection:  $\Pi$
  - Cartesian Product:  $\times$
- Derived or auxiliary operators:
  - Intersection, complement
  - Joins (natural, equi-join, theta join, semi-join)
  - Renaming:  $\rho$

# The Role of Relational Algebra in a DBMS



# 1. Union and 2. Difference

- $R1 \cup R2$
- Example:
  - $\text{ActiveEmployees} \cup \text{RetiredEmployees}$
- $R1 - R2$
- Example:
  - $\text{AllEmployees} - \text{RetiredEmployees}$

# What about Intersection ?

- It is a derived operator
- $R1 \cap R2 = R1 - (R1 - R2)$
- Also expressed as a join (will see later)
- Example
  - `UnionizedEmployees`  $\cap$  `RetiredEmployees`



# Union Compatible Relations

- Two relations are *union compatible* if
  - Both have same number of columns
  - Names of attributes are the same in both
  - Attributes with the same name in both relations have the same domain
- Union compatible relations can be combined using *union*, *intersection*, and *set difference*

# Example

Tables:

Person (*SSN, Name, Address, Hobby*)

Professor (*Id, Name, Office, Phone*)

are not union compatible.

But

$\pi_{Name}(\text{Person})$  and  $\pi_{Name}(\text{Professor})$   
are union compatible so

$\pi_{Name}(\text{Person}) - \pi_{Name}(\text{Professor})$   
makes sense.

# 3. Selection

- Returns all tuples which satisfy a condition
- Notation:  $\sigma_c(R)$
- Examples
  - $\sigma_{\text{Salary} > 40000}(\text{Employee})$
  - $\sigma_{\text{name} = \text{"Smith"}}(\text{Employee})$
- The condition  $c$  can be  $=, <, \leq, >, \geq, <>$

# Select Operator

- Produce table containing subset of rows of argument table satisfying condition

$$\sigma_{condition}(relation)$$

- Example:

Person

<i>Id</i>	<i>Name</i>	<i>Address</i>	<i>Hobby</i>
1123	John	123 Main	stamps
1123	John	123 Main	coins
5556	Mary	7 Lake Dr	hiking
9876	Bart	5 Pine St	stamps

$\sigma_{Hobby='stamps'}(Person)$

<i>Id</i>	<i>Name</i>	<i>Address</i>	<i>Hobby</i>
1123	John	123 Main	stamps
9876	Bart	5 Pine St	stamps

# Selection Condition

- Operators:  $<$ ,  $\leq$ ,  $\geq$ ,  $>$ ,  $=$ ,  $\neq$
- Simple selection condition:
  - *$\langle attribute \rangle operator \langle constant \rangle$*
  - *$\langle attribute \rangle operator \langle attribute \rangle$*
- *$\langle condition \rangle AND \langle condition \rangle$*
- *$\langle condition \rangle OR \langle condition \rangle$*
- NOT  *$\langle condition \rangle$*

# Selection Condition - Examples

- $\sigma_{Id > 3000 \text{ OR } Hobby = 'hiking'} (Person)$
- $\sigma_{Id > 3000 \text{ AND } Id < 3999} (Person)$
- $\sigma_{NOT(Hobby = 'hiking')} (Person)$
- $\sigma_{Hobby \neq 'hiking'} (Person)$

SSN	Name	Salary
1234545	John	200000
5423341	Smith	600000
4352342	Fred	500000

$\sigma_{\text{Salary} > 40000}$  (Employee)

SSN	Name	Salary
5423341	Smith	600000
4352342	Fred	500000

## 4. Projection

- Eliminates columns, then removes duplicates
- Notation:  $\Pi_{A1, \dots, An}(R)$
- Example: project social-security number and names:
  - $\Pi_{SSN, Name}(Employee)$
  - Output schema: Answer(SSN, Name)



SSN	Name	Salary
1234545	John	200000
5423341	John	600000
4352342	John	200000

$\Pi_{\text{Name,Salary}}(\text{Employee})$

Name	Salary
John	20000
John	60000

# Project Operator

- Produces table containing subset of columns of argument table

$$\pi_{\text{attribute list}}(\text{relation})$$

- Example:

Person

<i>Id</i>	<i>Name</i>	<i>Address</i>	<i>Hobby</i>
1123	John	123 Main	stamps
1123	John	123 Main	coins
5556	Mary	7 Lake Dr	hiking
9876	Bart	5 Pine St	stamps

$\pi_{\text{Name, Hobby}}(\text{Person})$

<i>Name</i>	<i>Hobby</i>
John	stamps
John	coins
Mary	hiking
Bart	stamps

# Project Operator

- Example:

Person

<i>Id</i>	<i>Name</i>	<i>Address</i>	<i>Hobby</i>
1123	John	123 Main	stamps
1123	John	123 Main	coins
5556	Mary	7 Lake Dr	hiking
9876	Bart	5 Pine St	stamps

$\pi_{Name,Address}(\text{Person})$

<i>Name</i>	<i>Address</i>
John	123 Main
Mary	7 Lake Dr
Bart	5 Pine St

Result is a table (no duplicates); can have fewer tuples than the original

# Expressions

$\pi_{Id, Name} (\sigma_{Hobby='stamps' \text{ OR } Hobby='coins'} (Person) )$

<i>Id</i>	<i>Name</i>	<i>Address</i>	<i>Hobby</i>
1123	John	123 Main	stamps
1123	John	123 Main	coins
5556	Mary	7 Lake Dr	hiking
9876	Bart	5 Pine St	stamps

Person

<i>Id</i>	<i>Name</i>
1123	John
9876	Bart

Result

## 5. Cartesian Product

- Each tuple in R1 with each tuple in R2
- Notation:  $R1 \times R2$
- Example:
  - Employee  $\times$  Dependents
- Very rare in practice; mainly used to express joins

## Cartesian Product Example

### Employee

Name	SSN
John	999999999
Tony	777777777

### Dependents

EmployeeSSN	Dname
999999999	Emily
777777777	Joe

### Employee x Dependents

Name	SSN	EmployeeSSN	Dname
John	999999999	999999999	Emily
John	999999999	777777777	Joe
Tony	777777777	999999999	Emily
Tony	777777777	777777777	Joe

# Cartesian Product

- If  $R$  and  $S$  are two relations,  $R \times S$  is the set of all concatenated tuples  $\langle x, y \rangle$ , where  $x$  is a tuple in  $R$  and  $y$  is a tuple in  $S$ 
  - $R$  and  $S$  need not be union compatible.
  - *But  $R$  and  $S$  must have distinct attribute names. Why?*
- $R \times S$  is expensive to compute. But why?

$A$	$B$	$C$	$D$
x1	x2	y1	y2
x3	x4	y3	y4

$R$

$A$	$B$	$C$	$D$
y1	y2	x1	x2
y3	y4	x3	x4

$S$

$A$	$B$	$C$	$D$
x1	x2	y1	y2
x1	x2	y3	y4
x3	x4	y1	y2
x3	x4	y3	y4

$R \times S$

# Relational Algebra

- Five operators:
  - Union:  $\cup$
  - Difference:  $-$
  - Selection:  $\sigma$
  - Projection:  $\Pi$
  - Cartesian Product:  $\times$
- Derived or auxiliary operators:
  - Intersection, complement
  - Joins (natural, equi-join, theta join, semi-join)
  - Renaming:  $\rho$



# Renaming

- Changes the schema, not the instance
- Notation:  $\rho_{B_1, \dots, B_n}(R)$
- Example:
  - $\rho_{\text{LastName}, \text{SocSocNo}}(\text{Employee})$
  - Output schema:  
Answer(LastName, SocSocNo)

# Renaming Example

## Employee

Name	SSN
John	999999999
Tony	777777777

## $\rho_{LastName, SocSocNo}$ (**Employee**)

LastName	SocSocNo
John	999999999
Tony	777777777

# Example – Second Method

Transcript (*StudId*, *CrsCode*, *Semester*, *Grade*)

Teaching (*ProfId*, *CrsCode*, *Semester*)

$$\pi_{StudId, CrsCode}(\text{Transcript})[StudId, CrsCode1]$$
$$\times \pi_{ProfId, CrsCode}(\text{Teaching})[ProfId, CrsCode2]$$

This is a relation with 4 attributes:

*StudId*, *CrsCode1*, *ProfId*, *CrsCode2*

# Natural Join

- Notation:  $R1 \bowtie R2$
- Meaning:  $R1 \bowtie R2 = \Pi_A(\sigma_C(R1 \times R2))$
- Where:
  - The selection  $\sigma_C$  checks equality of all common attributes
  - The projection eliminates the duplicate common attributes

## Natural Join Example

### Employee

Name	SSN
John	999999999
Tony	777777777

### Dependents

SSN	Dname
999999999	Emily
777777777	Joe

**Employee** ⋈ **Dependents** =

$\Pi_{\text{Name, SSN, Dname}}(\sigma_{\text{SSN}=\text{SSN}_2}(\text{Employee} \times \rho_{\text{SSN}_2, \text{Dname}}(\text{Dependents})))$

Name	SSN	Dname
John	999999999	Emily
Tony	777777777	Joe

# Theta Join

- A join that involves a predicate
- $R1 \bowtie_{\theta} R2 = \sigma_{\theta}(R1 \times R2)$
- Here  $\theta$  can be any condition

# Eq-join

- A theta join where  $\theta$  is an equality
- $R1 \mid \times \mid_{A=B} R2 = \sigma_{A=B} (R1 \times R2)$
- Example:
  - $\text{Employee} \mid \times \mid_{SSN=SSN} \text{Dependents}$
- Most useful join in practice

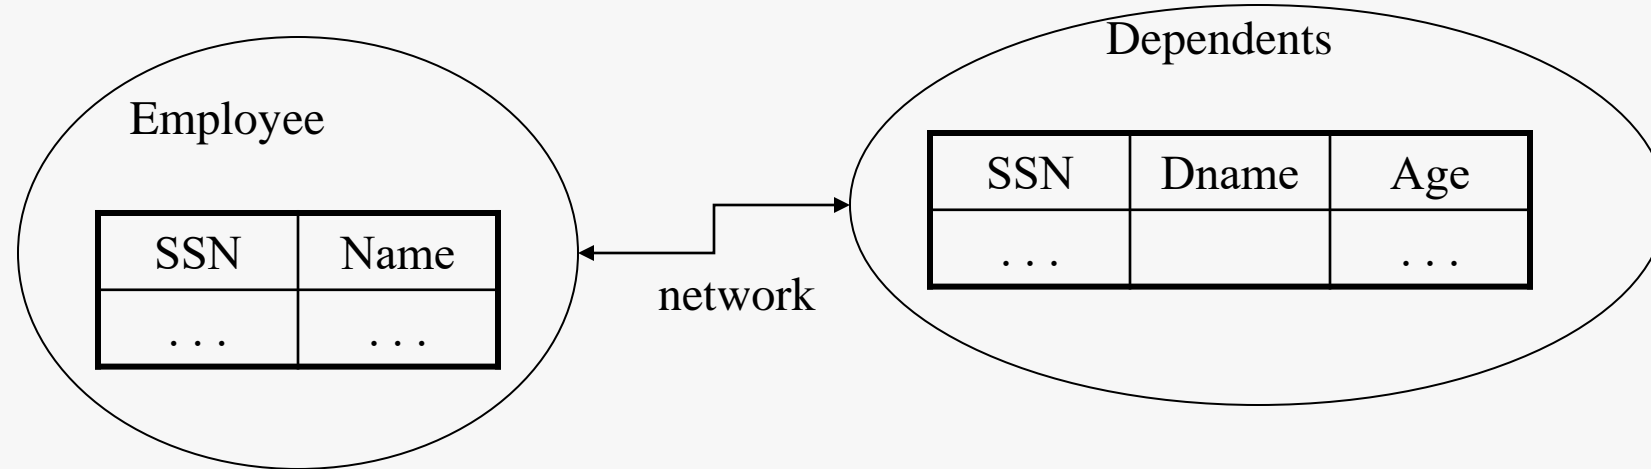
# Semijoin

- $R \bowtie S = \Pi_{A_1, \dots, A_n} (R \bowtie S)$
- Where  $A_1, \dots, A_n$  are the attributes in  $R$
- Example:
  - Employee  $\bowtie$  Dependents



# Semijoins in Distributed Databases

- Semijoins are used in distributed databases



$\text{Employee} \mid \times_{\text{ssn}=\text{ssn}} (\sigma_{\text{age}>71} (\text{Dependents}))$

$R = \text{Employee} \mid \times T$

$T = \Pi_{\text{SSN}} \sigma_{\text{age}>71} (\text{Dependents})$

$\text{Answer} = R \mid \times \mid \text{Dependents}$

# -- Aggregate Functions and Operations

- **Aggregation function** takes a collection of values and returns a single value as a result.

**avg:** average value

**min:** minimum value

**max:** maximum value

**sum:** sum of values

**count:** number of values

- **Aggregate operation** in relational algebra

$$G_1, G_2, \dots, G_n \mathcal{G} F_1(A_1), F_2(A_2), \dots, F_n(A_n) (E)$$

- $E$  is any relational-algebra expression
- $G_1, G_2, \dots, G_n$  is a list of attributes on which to group (can be empty)
- Each  $F_i$  is an aggregate function
- Each  $A_i$  is an attribute name

# --- Aggregate Operation – Example 1

<i>A</i>	<i>B</i>	<i>C</i>
$\alpha$	$\alpha$	7
$\alpha$	$\beta$	7
$\beta$	$\beta$	3
$\beta$	$\beta$	10

***r***

$\Pi_{g^{\text{sum}(c)}(\mathbf{r})} \longrightarrow$

<i>sum-C</i>
27

## --- Aggregate Operation – Example 2

- Relation *account* grouped by *branch-name*:

<i>branch-name</i>	<i>account-number</i>	<i>balance</i>
Dammam	A-102	400
Dammam	A-201	900
Khobar	A-217	750
Khobar	A-215	750
Hafuf	A-222	700

$T = \text{branch-name } \mathcal{G}_{\text{sum}(\text{balance})}(\text{account})$

<i>branch-name</i>	<i>balance</i>
Dammam	1300
Khobar	1500
Hafuf	700

## --- Aggregate Functions: Renaming

- Result of aggregation does not have a name
  - Can use rename operation to give it a name
  - For convenience, we permit renaming as part of aggregate operation

*branch-name*  $\mathcal{G}$  *sum(balance) as sum-balance* (*account*)

# Finally: RA has Limitations !

- Cannot compute “transitive closure”

Name1	Name2	Relationship
Fred	Mary	Father
Mary	Joe	Cousin
Mary	Bill	Spouse
Nancy	Lou	Sister

- Find all direct and indirect relatives of Fred
- Cannot express in RA !!! Need to write C program

# SQL to Relational Algebra Conversation

# Schema for Student Registration System

Student (Id, Name, Addr, Status)

Professor (Id, Name, DeptId)

Course (DeptId, CrsCode, CrsName, Descr)

Transcript (StudId, CrsCode, Semester, Grade)

Teaching (ProfId, CrsCode, Semester)

Department (DeptId, Name)



# Query Sublanguage of SQL

```
SELECT  C.CrsName  
FROM    Course C  
WHERE   C.DeptId = 'CS'
```

- *Tuple variable* C ranges over rows of Course.
- Evaluation strategy:
  - FROM clause produces Cartesian product of listed tables
  - WHERE clause assigns rows to C in sequence and produces table containing only rows satisfying condition
  - SELECT clause retains listed columns
- Equivalent to:  $\pi_{CrName} \sigma_{DeptId='CS'}(\text{Course})$

# Join Queries

```
SELECT C.CrsName  
FROM Course C, Teaching T  
WHERE C.CrsCode=T.CrsCode AND T.Semester='S2000'
```

- List CS courses taught in S2000
- Tuple variables clarify meaning.
- Join condition “*C.CrsCode=T.CrsCode*”
  - relates facts to each other
- Selection condition “ *T.Semester='S2000'* ”
  - eliminates irrelevant rows
- Equivalent (using natural join) to:

$$\pi_{CrName}(\text{Course} \bowtie \sigma_{Semester='S2000'}(\text{Teaching}))$$
$$\pi_{CrName}(\sigma_{Sem='S2000'}(\text{Course} \bowtie \text{Teaching}))$$

# Relational Algebra and SQL Exercises

- Professor(ssn, profname, status)
- Course(crscode, crsname, credits)
- Taught(crscode, semester, ssn)

# Query 1

```
Professor(ssn, profname, status)  
Course(crscode, crsname, credits)  
Taught(crscode, semester, ssn)
```

Return those professors who have  
taught 'csc6710' but never 'csc7710'.

# Relational Algebra Solution

Professor(ssn, profname, status)  
Course(crscode, crsname, credits)  
Taught(crscode, semester, ssn)

$\pi_{\text{ssn}}(\sigma_{\text{crscode}='csc6710'}(\text{Taught})) -$   
 $\pi_{\text{ssn}}(\sigma_{\text{crscode}='csc7710'}(\text{Taught}))$

# SQL Solution

```
Professor(ssn, profname, status)
Course(crscode, crsname, credits)
Taught(crscode, semester, ssn)
```

```
(SELECT ssn
From Taught
Where crscode = 'CSC6710')
EXCEPT
(SELECT ssn
From Taught
Where crscode = 'CSC7710'))
```

## Query 2

```
Professor(ssn, profname, status)  
Course(crscode, crsname, credits)  
Taught(crscode, semester, ssn)
```

Return those professors who have  
taught both ‘csc6710’ and ‘csc7710’.

# Relational Algebra Solution

Professor(ssn, profname, status)  
Course(crscode, crsname, credits)  
Taught(crscode, semester, ssn)

$\pi_{\text{ssn}}(\sigma_{\text{crscode}='csc6710' \wedge \text{crscode}='csc7710'}(\text{Taught}))$ ,  
wrong!

$\pi_{\text{ssn}}(\sigma_{\text{crscode}='csc6710'}(\text{Taught})) \cap$   
 $\pi_{\text{ssn}}(\sigma_{\text{crscode}='csc7710'}(\text{Taught}))$ , correct!



# SQL Solution

```
Professor(ssn, profname, status)
Course(crscode, crsname, credits)
Taught(crscode, semester, ssn)
```

```
SELECT T1.ssn
From Taught T1, Taught T2,
Where T1.crscode = 'CSC6710' AND T2.crscode='CSC7710' AND
T1.ssn=T2.ssn
```

## Query 3

```
Professor(ssn, profname, status)  
Course(crscode, crsname, credits)  
Taught(crscode, semester, ssn)
```

Return those professors who have never taught 'csc7710'.

# Relational Algebra Solution

Professor(ssn, profname, status)  
Course(crscode, crsname, credits)  
Taught(crscode, semester, ssn)

$\pi_{\text{ssn}}(\sigma_{\text{crscode} \neq \text{'csc7710'}}(\text{Taught}))$ ,

wrong answer!

$\pi_{\text{ssn}}(\text{Professor}) - \pi_{\text{ssn}}(\sigma_{\text{crscode} = \text{'csc7710'}}(\text{Taught}))$ ,

correct answer!

# SQL Solution

```
Professor(ssn, profname, status)  
Course(crscode, crsname, credits)  
Taught(crscode, semester, ssn)
```

```
(SELECT ssn  
From Professor)  
EXCEPT  
(SELECT ssn  
From Taught T  
Where T.crscode = 'CSC7710')
```

## Query 4

```
Professor(ssn, profname, status)  
Course(crscode, crsname, credits)  
Taught(crscode, semester, ssn)
```

Return those professors who taught  
‘CSC6710’ and ‘CSC7710’ in the same  
semester

# Relational Algebra Solution

Professor(ssn, profname, status)

Course(crscode, crsname, credits)

Taught(crscode, semester, ssn)

## Relational Algebra Solution

$\pi_{ssn}(\sigma_{crscode1='csc6710'}(Taught[crscode1, ssn, semester])) \bowtie$

$\sigma_{crscode2='csc7710'}(Taught[crscode2, ssn, semester]))$

# SQL Solution

```
Professor(ssn, profname, status)
Course(crscode, crsname, credits)
Taught(crscode, semester, ssn)
```

```
SELECT T1.ssn
From Taught T1, Taught T2,
Where T1.crscode = 'CSC6710' AND T2.crscode='CSC7710' AND
T1.ssn=T2.ssn AND T1.semester=T2.semester
```

## Query 5

```
Professor(ssn, profname, status)  
Course(crscode, crsname, credits)  
Taught(crscode, semester, ssn)
```

Return those professors who taught  
‘CSC6710’ or ‘CSC7710’ but not both.



# Relational Algebra Solution

Professor(ssn, profname, status)  
Course(crscode, crsname, credits)  
Taught(crscode, semester, ssn)

$$\pi_{\text{ssn}}(\sigma_{\text{crscode} \neq \text{'csc7710'} \vee \text{crscode} = \text{'csc7710'}, (\text{Taught})) -$$
$$(\pi_{\text{ssn}}(\sigma_{\text{crscode} = \text{'csc6710'}, (\text{Taught}))) \cap$$
$$\pi_{\text{ssn}}(\sigma_{\text{crscode} = \text{'csc7710'}, (\text{Taught})))$$

# SQL Solution

Professor(ssn, profname, status)

Course(crscode, crsname, credits)

Taught(crscode, semester, ssn)

```
(SELECT ssn
FROM Taught T
WHERE T.crscode='CSC6710' OR T.crscode='CSC7710')
Except
(SELECT T1.ssn
From Taught T1, Taught T2,
Where T1.crscode = 'CSC6710') AND T2.crscode='CSC7710' AND
T1.ssn=T2.ssn)
```

## Query 6

```
Professor(ssn, profname, status)  
Course(crscode, crsname, credits)  
Taught(crscode, semester, ssn)
```

Return those courses that have never been taught.

# Relational Algebra Solution

Professor(ssn, profname, status)

Course(crscode, crsname, credits)

Taught(crscode, semester, ssn)

$$\pi_{\text{crscode}}(\text{Course}) - \pi_{\text{crscode}}(\text{Taught})$$

# SQL Solution

```
Professor(ssn, profname, status)  
Course(crscode, crsname, credits)  
Taught(crscode, semester, ssn)
```

```
(SELECT crscode  
FROM Course)  
EXCEPT  
(SELECT crscode  
FROM TAUGHT  
)
```

# Query 7

```
Professor(ssn, profname, status)  
Course(crscode, crsname, credits)  
Taught(crscode, semester, ssn)
```

Return those courses that have been taught at least in two semesters.

# Relational Algebra Solution

Professor(ssn, profname, status)  
Course(crscode, crsname, credits)  
Taught(crscode, semester, ssn)

$\pi_{\text{crscode}}(\sigma_{\text{semester1} \neq \text{semester2}}(\text{Taught}[\text{crscode}, \text{ssn1}, \text{semester1}] \bowtie \text{Taught}[\text{crscode}, \text{ssn2}, \text{semester2}])))$

# SQL Solution

```
Professor(ssn, profname, status)  
Course(crscode, crsname, credits)  
Taught(crscode, semester, ssn)
```

```
SELECT T1.crscode  
FROM Taught T1, Taught T2  
WHERE T1.crscode=T2.crscode AND T1.semester <> T2.semester
```



## Query 8

```
Professor(ssn, profname, status)  
Course(crscode, crsname, credits)  
Taught(crscode, semester, ssn)
```

Return those courses that have been taught at least in 10 semesters.

# SQL Solution

Professor(ssn, profname, status)

Course(crscode, crsname, credits)

Taught(crscode, semester, ssn)

```
SELECT crscode  
FROM Taught  
GROUP BY crscode  
HAVING COUNT(*) >= 10
```

## Query 9

```
Professor(ssn, profname, status)  
Course(crscode, crsname, credits)  
Taught(crscode, semester, ssn)
```

Return those courses that have been taught by at least 5 different professors.

# SQL Solution

Professor(ssn, profname, status)

Course(crscode, crsname, credits)

Taught(crscode, semester, ssn)

```
SELECT crscode
FROM (SELECT DISTINCT crscode, ssn FROM TAUGHT)
GROUP BY crscode
HAVING COUNT(*) >= 5
```

# Query 10

```
Professor(ssn, profname, status)  
Course(crscode, crsname, credits)  
Taught(crscode, semester, ssn)
```

Return the names of professors who  
ever taught 'CSC6710'.

# Relational Algebra Solution

Professor(ssn, profname, status)  
Course(crscode, crsname, credits)  
Taught(crscode, semester, ssn)

$\pi_{\text{profname}}(\sigma_{\text{crscode}='csc6710'}(\text{Taught}) \bowtie \text{Professor})$

# SQL Solution

```
Professor(ssn, profname, status)  
Course(crscode, crsname, credits)  
Taught(crscode, semester, ssn)
```

```
SELECT P.profname  
FROM Professor P, Taught T  
WHERE P.ssn = T.ssn AND T.crscode = 'CSC6710'
```

# Query 11

```
Professor(ssn, profname, status)  
Course(crscode, crsname, credits)  
Taught(crscode, semester, ssn)
```

Return the names of full professors  
who ever taught 'CSC6710'.



# Relational Algebra Solution

Professor(ssn, profname, status)  
Course(crscode, crsname, credits)  
Taught(crscode, semester, ssn)

$\pi_{\text{profname}}(\sigma_{\text{crscode}='csc6710'}(\text{Taught}) \bowtie$   
 $\sigma_{\text{status}='full'}(\text{Professor}))$

# SQL Solution

Professor(ssn, profname, status)

Course(crscode, crsname, credits)

Taught(crscode, semester, ssn)

```
SELECT P.profname  
FROM Professor P, Taught T  
WHERE P.status = 'full' AND P.ssn = T.ssn AND T.crscode =  
'CSC6710'
```

# Query 12

```
Professor(ssn, profname, status)  
Course(crscode, crsname, credits)  
Taught(crscode, semester, ssn)
```

Return the names of full professors  
who ever taught more than two courses  
in one semester.

# SQL Solution

Professor(ssn, profname, status)

Course(crscode, crsname, credits)

Taught(crscode, semester, ssn)

```
SELECT P.profname
FROM Professor P
WHERE ssn IN(
SELECT ssn
FROM Taught
GROUP BY ssn, semester
HAVING COUNT(*) > 2
)
```

# Query 13

```
Professor(ssn, profname, status)  
Course(crscode, crsname, credits)  
Taught(crscode, semester, ssn)
```

Delete those professors who never taught a course.

# SQL Solution

```
Professor(ssn, profname, status)  
Course(crscode, crsname, credits)  
Taught(crscode, semester, ssn)
```

```
DELETE FROM Professor  
WHERE ssn NOT IN  
(SELECT ssn  
FROM Taught  
)
```

# Query 14

```
Professor(ssn, profname, status)  
Course(crscode, crsname, credits)  
Taught(crscode, semester, ssn)
```

Change all the credits to 4 for those courses that are taught in f2006 semester.

# SQL Solution

Professor(ssn, profname, status)

Course(crscode, crsname, credits)

Taught(crscode, semester, ssn)

```
UPDATE Course
SET credits = 4
WHERE crscode IN
(
    SELECT crscode
    FROM Taught
    WHERE semester = 'f2006'
)
```



# Query 15

```
Professor(ssn, profname, status)  
Course(crscode, crsname, credits)  
Taught(crscode, semester, ssn)
```

Return the names of the professors who have taught more than 30 credits of courses.

# SQL Solution

```
Professor(ssn, profname, status)
Course(crscode, crsname, credits)
Taught(crscode, semester, ssn)
```

```
SELECT profname
FROM Professor
WHERE ssn IN
(
    SELECT T.ssn
    FROM Taught T, Course C
    WHERE T.crscode = C.crscode
    GROUP BY T.ssn
    HAVING SUM(C.credits) > 30
)
```

# Query 16

```
Professor(ssn, profname, status)  
Course(crscode, crsname, credits)  
Taught(crscode, semester, ssn)
```

Return the name(s) of the professor(s) who taught the most number of courses in S2006.

# SQL Solution

Professor(ssn, profname, status)

Course(crscode, crsname, credits)

Taught(crscode, semester, ssn)

```
SELECT profname
FROM Professor
WHERE ssn IN(
    SELECT ssn FROM Taught
    WHERE semester = 'S2006'
    GROUP BY ssn
    HAVING COUNT(*) =
        (SELECT MAX(Num)
         FROM
            (SELECT ssn, COUNT(*) as Num
             FROM Taught
             WHERE semester = 'S2006'
             GROUP BY ssn)
         )
)
```

# Query 17

```
Professor(ssn, profname, status)  
Course(crscode, crsname, credits)  
Taught(crscode, semester, ssn)
```

List all the course names that professor  
‘Smith’ taught in Fall of 2007.

# Relational Algebra Solution

Professor(ssn, profname, status)  
Course(crscode, crsname, credits)  
Taught(crscode, semester, ssn)

$\pi_{\text{crsname}}(\sigma_{\text{profname}='Smith'}(\text{Professor}) \bowtie$   
 $\sigma_{\text{semester}='f2007'}(\text{Taught}) \bowtie$   
 $\text{Course})$

# SQL Solution

```
Professor(ssn, profname, status)  
Course(crscode, crsname, credits)  
Taught(crscode, semester, ssn)
```

```
SELECT crsname  
FROM Professor P, Taught T, Course C  
WHERE P.profname = 'Smith' AND P.ssn = T.ssn AND  
T.semester = 'F2007' AND T.crscode = C.crscode
```

# Query 18

```
Professor(ssn, profname, status)  
Course(crscode, crsname, credits)  
Taught(crscode, semester, ssn)
```

In chronological order, list the number of courses that the professor with ssn ssn = 123456789 taught in each semester.



# SQL Solution

Professor(ssn, profname, status)

Course(crscode, crsname, credits)

Taught(crscode, semester, ssn)

```
SELECT semester, COUNT(*)  
FROM Taught  
WHERE ssn = '123456789'  
GROUP BY semester  
ORDER BY semester ASC
```

# Query 19

```
Professor(ssn, profname, status)  
Course(crscode, crsname, credits)  
Taught(crscode, semester, ssn)
```

In alphabetical order of the names of professors, list the name of each professor and the total number of courses she/he has taught.

# SQL Solution

Professor(ssn, profname, status)

Course(crscode, crsname, credits)

Taught(crscode, semester, ssn)

```
SELECT P.profname, COUNT(*)  
FROM Professor P, Taught T  
WHERE P.ssn = T.ssn  
GROUP BY P.ssn, P.profname  
ORDER BY P.profname ASC
```

## Query 20

```
Professor(ssn, profname, status)  
Course(crscode, crsname, credits)  
Taught(crscode, semester, ssn)
```

Delete those professors who taught less than 10 courses.

# SQL Solution

Professor(ssn, profname, status)

Course(crscode, crsname, credits)

Taught(crscode, semester, ssn)

```
DELETE FROM Professor
WHERE ssn IN(
    SELECT ssn
    FROM Taught
    GROUP BY ssn
    HAVING COUNT(*) < 10
)
```

# Query 21

```
Professor(ssn, profname, status)  
Course(crscode, crsname, credits)  
Taught(crscode, semester, ssn)
```

Delete those professors who taught less than 40 credits.

# SQL Solution

```
Professor(ssn, profname, status)
Course(crscode, crsname, credits)
Taught(crscode, semester, ssn)
```

```
DELETE FROM Professor
WHERE ssn IN(
    SELECT T.ssn
    FROM Taught T, Course C
    WHERE T.crscode = C.crscode
    GROUP BY ssn
    HAVING SUM(C.credits) < 40
)
```

## Query 22

```
Professor(ssn, profname, status)  
Course(crscode, crsname, credits)  
Taught(crscode, semester, ssn)
```

List those professors who have not taught any course in the past three semesters (F2006, W2007, F2007).



# SQL Solution

```
Professor(ssn, profname, status)
Course(crscode, crsname, credits)
Taught(crscode, semester, ssn)
```

```
SELECT *
FROM Professor P
WHERE NOT EXISTS(
    SELECT *
    FROM Taught
    WHERE P.ssn = T.ssn AND (T.semester = 'F2006' OR
    T.semester = 'W2007' OR T.semester='F2007'))
)
```

## Query 23

```
Professor(ssn, profname, status)  
Course(crscode, crsname, credits)  
Taught(crscode, semester, ssn)
```

List the names of those courses that professor Smith have never taught.

# Relational Algebra Solution

Professor(ssn, profname, status)  
Course(crscode, crsname, credits)  
Taught(crscode, semester, ssn)

$\pi_{\text{crsname}}(\text{Course}) -$

$\pi_{\text{crsname}}(\sigma_{\text{profname}='Smith'}(\text{Professor}) \bowtie$   
 $(\text{Taught}) \bowtie$

Course)

# SQL Solution

```
Professor(ssn, profname, status)
Course(crscode, crsname, credits)
Taught(crscode, semester, ssn)
```

```
SELECT crsname
FROM Course C
WHERE NOT EXISTS
    SELECT *
    FROM Professor P, Taught T
    WHERE P.profname='Smith' AND P.ssn = T.ssn AND
    T.crscode = C.crscode
)
```

## Query 24

```
Professor(ssn, profname, status)  
Course(crscode, crsname, credits)  
Taught(crscode, semester, ssn)
```

Return those courses that have been taught by all professors.

# Relational Algebra Solution

Professor(ssn, profname, status)  
Course(crscode, crsname, credits)  
Taught(crscode, semester, ssn)

$$\pi_{\text{crscode, ssn}}(\text{Taught}) / \pi_{\text{ssn}}(\text{Professor})$$

# SQL Solution

Professor(ssn, profname, status)

Course(crscode, crsname, credits)

Taught(crscode, semester, ssn)

```
SELECT crscode
FROM Taught T1
WHERE NOT EXISTS(
    (SELECT ssn
     FROM Professor)
    EXCEPT
    (SELECT ssn
     FROM Taught T2
     WHERE T2.crscode = T1.crscode)
)
```

## Query 25

```
Professor(ssn, profname, status)  
Course(crscode, crsname, credits)  
Taught(crscode, semester, ssn)
```

Return those courses that have been taught in all semesters.



# Relational Algebra Solution

Professor(ssn, profname, status)  
Course(crscode, crsname, credits)  
Taught(crscode, semester, ssn)

$$\pi_{\text{crscode, semester}}(\text{Taught}) / \pi_{\text{semester}}(\text{Taught})$$

# SQL Solution

Professor(ssn, profname, status)

Course(crscode, crsname, credits)

Taught(crscode, semester, ssn)

```
SELECT crscode
FROM Taught T1
WHERE NOT EXISTS(
    (SELECT semester
     FROM Taught)
    EXCEPT
    (SELECT semester
     FROM Taught T2
     WHERE T2.crscode = T1.crscode)
)
```

## Query 25

```
Professor(ssn, profname, status)  
Course(crscode, crsname, credits)  
Taught(crscode, semester, ssn)
```

Return those courses that have been taught **ONLY** by junior professors.

# Relational Algebra Solution

Professor(ssn, profname, status)  
Course(crscode, crsname, credits)  
Taught(crscode, semester, ssn)

$\pi_{\text{crscode}}(\text{Course}) - \pi_{\text{crscode}}$   
 $(\sigma_{\text{status} \neq \text{'Junior'}}(\text{Professor}) \bowtie \text{Taught})$

# SQL Solution

Professor(ssn, profname, status)

Course(crscode, crsname, credits)

Taught(crscode, semester, ssn)

```
SELECT crscode
FROM Course C
WHERE c.crscode NOT IN(
    (SELECT crscode
     FROM Taught T, Professor P
     WHERE T.ssn = P.ssn AND P.status='Junior'
    )
)
```



Thank  
You

A blue paper cutout of the words "Thank You" in a stylized, rounded font. The text is white with a blue outline. The cutout is hanging from a thin brown string that is tied in a loop at the top. The background is white.